

[Stack Overflow](#) | [Google Group](#) | [Gitter Chat](#) | [Subreddit](#) | [Youtube Channel](#) | [Documentation](#) | [Contribution Guide](#) |

Framework	Language	Max Throughput	Avg Latency	Transfer
Go FastHttp	Go	1,396,685.83	99.98ms	167.83MB
Light-4j	Java	1,344,512.65	2.36ms	169.25MB
ActFramework	Java	945,429.13	2.22ms	136.15MB
Go Iris	Go	828,035.66	5.77ms	112.92MB
Vertx	Java	803,311.31	2.37ms	98.06MB
Node-uws	Node/C++	589,924.44	7.22ms	28.69MB
Dotnet	.Net	486,216.93	2.93ms	57.03MB
Jooby/Undertow	Java	362,018.07	3.95ms	47.99MB
SeedStack-Filter	Java	343,416.33	4.41ms	51.42MB
Spring Boot Reactor	Java	243,240.17	7.44ms	17.86MB
Pippo-Undertow	Java	216,254.56	9.80ms	31.35MB
Spark	Java	194,553.83	13.85ms	32.47MB
Pippo-Jetty	Java	178,055.45	15.66ms	26.83MB
Play-Java	Java	177,202.75	12.15ms	21.80MB
Go HttpRouter	Go	171,852.31	14.12ms	21.14MB
Go Http	Go	170,313.02	15.01ms	20.95MB
JFinal	Java	139,467.87	11.89ms	29.79MB
Akka-Http	Java	132,157.96	12.21ms	19.54MB
RatPack	Java	124,700.70	13.45ms	10.82MB
Pippo-Tomcat	Java	103,948.18	23.50ms	15.29MB
Bootique + Jetty/Jersey	Java	65,072.20	39.08ms	11.17MB
SeedStack-Jersey2	Java	52,310.11	26.88ms	11.87MB
Baseio	Java	50,361.98	22.20ms	6.39MB
NinjaFramework	Java	47,956.43	55.76ms	13.67MB
Play 1	Java	44,491.87	10.73ms	18.75MB
Spring Boot Undertow	Java	44,260.61	38.94ms	6.42MB
Nodejs Express	Node	42,443.34	22.30ms	9.31MB
Dropwizard	Java	33,819.90	98.78ms	3.23MB

Framework	Language	Max Throughput	Avg Latency	Transfer
Spring Boot Tomcat	Java	33,086.22	82.93ms	3.98MB
Node-Loopback	Node	32,091.95	34.42ms	11.51MB
Payra Micro	Java	24,768.69	118.86ms	3.50MB
WildFly Swarm	Java	21,541.07	59.77ms	2.83MB

We are using pipeline.lua to generate more requests per second and the pipeline.lua is located at `microservices-framework-benchmark/pipeline.lua`.

Here is the light-java server performance with the same command line with other frameworks.

```
wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --latency -- / 16
Running 30s test @ http://localhost:8080
4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency   774.30us    1.77ms  40.31ms   91.51%
  Req/Sec    0.93M    81.46k   1.06M    87.83%
Latency Distribution
  50%    286.00us
  75%    529.00us
  90%     1.92ms
  99%     0.00us
110603088 requests in 30.01s, 12.88GB read
Requests/sec: 3685234.33
Transfer/sec:   439.31MB
```

Here is another test with for light-java to push more requests but some other frameworks get server errors.

```
steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 50
Running 30s test @ http://localhost:8080
4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency     2.35ms    2.80ms  54.23ms   89.24%
  Req/Sec   550.10k    64.34k   1.22M    74.54%
Latency Distribution
  50%     1.56ms
  75%     2.87ms
  90%     5.41ms
  99%    22.17ms
65803650 requests in 30.10s, 6.50GB read
Requests/sec: 2186203.44
Transfer/sec:   221.00MB
```

Here is the spring-boot-tomcat (tomcat embedded) performance.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua -
-latency -- / 16
Running 30s test @ http://localhost:8080
4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    82.93ms   108.77ms   1.58s    89.45%
  Req/Sec    8.40k      3.68k    22.19k    68.54%
Latency Distribution
  50%    45.66ms
  75%   101.59ms
  90%   197.72ms
  99%   542.87ms
995431 requests in 30.09s, 119.79MB read
Requests/sec: 33086.22
Transfer/sec: 3.98MB

```

Here is the spring-boot-undertow (undertow embedded) performance.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    38.94ms   39.29ms  456.82ms    89.28%
  Req/Sec   11.21k     4.97k   28.16k    68.14%
Latency Distribution
  50%    27.58ms
  75%    49.62ms
  90%    80.73ms
  99%   201.87ms
1331312 requests in 30.08s, 192.98MB read
Requests/sec: 44260.61
Transfer/sec: 6.42MB

```

Basically, light-4j is 44 times faster than spring-boot with tomcat embedded just for the raw performance to serve Hello World!

In order to have a closer comparison, I have created another project spring-boot-undertow with embedded undertow servlet container (light-java is using undertow core only) and the performance is getting a little better. Light-Java is about 33 times faster than spring-boot with undertow embedded.

Upon requests from the community, I have added nodejs and golang examples and here are the testing result.

Node express framework. To start the server

```
cd node-express
npm install
node server.js
```

The test result.

```
steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
 4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    22.30ms   24.35ms  592.24ms   49.18%
  Req/Sec    10.70k     0.87k   11.95k    94.82%
Latency Distribution
  50%    47.94ms
  75%     0.00us
  90%     0.00us
  99%     0.00us
1274289 requests in 30.02s, 279.51MB read
Requests/sec: 42443.34
Transfer/sec:      9.31MB
```

Go Standard Http

To start the server

```
cd go-http
go run server.go -prefork
```

The testing result.

```
steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
 4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    15.01ms   15.35ms  180.11ms   87.10%
  Req/Sec    42.80k     5.46k   62.49k    70.58%
Latency Distribution
  50%    10.03ms
  75%    19.96ms
  90%    34.55ms
  99%    72.99ms
5123194 requests in 30.08s, 630.28MB read
```

```
Requests/sec: 170313.02
Transfer/sec:   20.95MB
```

## Go FastHttp

To start the server

```
cd go-fasthttp
go run server.go -prefork
```

The testing result.

```
steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
 4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    99.98ms  127.12ms 653.72ms   82.12%
  Req/Sec   351.24k    46.23k 525.74k    77.09%
Latency Distribution
 50%    30.76ms
 75%   175.44ms
 90%   299.14ms
 99%   476.20ms
41989168 requests in 30.06s, 4.93GB read
Requests/sec: 1396685.83
Transfer/sec:   167.83MB
```

After I post this online, one of spring developers recommended to test against Spring Boot with Reactor which is Netty based without servlet container. I am very new to this and might miss something and everyone is welcomed to submit pull request to enhance this project.

Here is the test result.

```
steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:3000 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:3000
 4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    7.44ms   12.88ms 285.71ms   94.23%
  Req/Sec   61.44k    12.25k  88.29k    79.23%
Latency Distribution
 50%     4.62ms
 75%     8.11ms
```

```

    90%    15.03ms
    99%    42.60ms
7305649 requests in 30.03s, 536.48MB read
Requests/sec: 243240.17
Transfer/sec:      17.86MB

```

Add Spark test case and here is the result. It is much better than most frameworks with servlet containers.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:4567 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:4567
4 threads and 128 connections
Thread Stats   Avg      Stdev     Max   +/-  Stdev
  Latency    13.85ms   30.58ms  687.89ms   96.29%
  Req/Sec    49.10k    16.55k   107.63k    73.41%
Latency Distribution
  50%     7.28ms
  75%    13.96ms
  90%    24.71ms
  99%   158.21ms
5855187 requests in 30.10s, 0.95GB read
Requests/sec: 194553.83
Transfer/sec:      32.47MB

```

Add Jooby test case and here is the result. It is better than Spring Boot as it is using Netty directly.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
4 threads and 128 connections
Thread Stats   Avg      Stdev     Max   +/-  Stdev
  Latency    15.57ms   25.97ms  431.00ms   94.64%
  Req/Sec    30.50k     8.66k    77.57k    75.68%
Latency Distribution
  50%     9.93ms
  75%    16.49ms
  90%    27.27ms
  99%   123.80ms
3613008 requests in 30.05s, 392.80MB read
Requests/sec: 120232.86
Transfer/sec:      13.07MB

```

@jknack submitted a pull request for Jooby to switch to Undertow instead of Netty and here is the updated result.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
 4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    13.22ms   20.07ms  439.26ms   94.77%
  Req/Sec    32.96k     7.71k   52.00k    78.38%
Latency Distribution
  50%     9.06ms
  75%    14.98ms
  90%    23.88ms
  99%    89.30ms
3926256 requests in 30.10s, 565.40MB read
Requests/sec: 130458.14
Transfer/sec:   18.79MB

```

@windbender submitted a pull request for Dropwizard and here is the result.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
 4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    98.78ms  167.67ms   1.92s   92.91%
  Req/Sec     8.58k    4.26k   24.02k   66.49%
Latency Distribution
  50%    47.95ms
  75%   103.68ms
  90%   211.24ms
  99%   958.22ms
1017904 requests in 30.10s, 97.07MB read
Requests/sec: 33819.90
Transfer/sec:   3.23MB

```

@IRus submitted a pull request for Bootique + Jetty/Jersey and here is the result.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
 4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    39.08ms   43.35ms  462.13ms   87.33%
  Req/Sec    16.46k     3.79k   31.04k    70.42%
Latency Distribution
  50%    23.94ms
  75%    49.95ms
  90%    94.78ms

```

```

    99% 204.74ms
1956621 requests in 30.07s, 335.88MB read
Requests/sec: 65072.20
Transfer/sec: 11.17MB

```

@hydra1983 submitted a pull request for RatPack and here is the result.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:5050 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:5050
4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
Latency       13.45ms   17.27ms  366.84ms   94.60%
Req/Sec       31.44k    6.39k   48.42k    83.28%
Latency Distribution
 50%    10.04ms
 75%    16.67ms
 90%    24.47ms
 99%    82.36ms
3747731 requests in 30.05s, 325.24MB read
Requests/sec: 124700.70
Transfer/sec: 10.82MB

```

@IRus submitted a pull request for WildFly Swarm and here is the result.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
Latency       59.77ms   44.97ms  526.38ms   64.92%
Req/Sec        5.41k    1.43k    9.53k    74.58%
Latency Distribution
 50%    74.16ms
 75%   162.88ms
 90%     0.00us
 99%     0.00us
648138 requests in 30.09s, 85.30MB read
Requests/sec: 21541.07
Transfer/sec: 2.83MB

```

@IRus submitted a pull request for Payra-Micro and here is the result.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080

```



```

4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency   118.86ms  144.34ms   1.77s    88.56%
  Req/Sec    6.29k     2.75k   15.40k    63.00%
Latency Distribution
  50%    69.49ms
  75%   155.56ms
  90%   282.09ms
  99%   696.93ms
745507 requests in 30.10s, 105.28MB read
Socket errors: connect 0, read 0, write 0, timeout 1
Requests/sec: 24768.69
Transfer/sec:    3.50MB

```

@hydra1983 submitted a pull request for akka-http and here is the result.

@cortwave submitted a pull request to upgrade to version 2.4.11. Here is the latest result.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    12.21ms   30.77ms  712.06ms   98.78%
  Req/Sec   33.52k     3.99k   68.36k    86.96%
Latency Distribution
  50%     8.57ms
  75%    13.30ms
  90%    18.62ms
  99%    52.13ms
3972310 requests in 30.06s, 587.18MB read
Requests/sec: 132157.96
Transfer/sec:    19.54MB

```

@dpapworth submitted a pull request to add play-java and here is the result.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua -
-latency -- / 16
Running 30s test @ http://localhost:8080
4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    12.15ms   14.27ms  375.72ms   89.98%
  Req/Sec   44.60k     7.32k   86.18k    73.83%
Latency Distribution
  50%     8.20ms
  75%    15.65ms
  90%    26.44ms

```

```

    99%    60.57ms
  5325273 requests in 30.05s, 655.14MB read
Requests/sec: 177202.75
Transfer/sec:    21.80MB

```

Add iris which is a go web framework based on fasthttp and here is the result.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
 4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    5.77ms   11.26ms  181.19ms   90.98%
 Req/Sec   208.26k   36.36k   347.64k    74.83%
Latency Distribution
 50%    2.06ms
 75%    4.54ms
 90%   15.21ms
 99%   57.44ms
24886288 requests in 30.05s, 3.31GB read
Requests/sec: 828035.66
Transfer/sec:    112.92MB

```

@greenlaw110 submitted a pull request to add actframework and here is the result.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
 4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    2.22ms    6.58ms  219.99ms   97.98%
 Req/Sec   241.95k   59.46k   420.20k    73.85%
Latency Distribution
 50%    1.45ms
 75%    2.24ms
 90%    3.65ms
 99%   11.87ms
28389984 requests in 30.03s, 3.99GB read
Requests/sec: 945429.13
Transfer/sec:    136.15MB

```

@greenlaw110 submitted a pull request to add JFinal and here is the result.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua -
-latency -- / 16
Running 30s test @ http://localhost:8080

```

```

4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    11.89ms   13.66ms  308.54ms   91.36%
  Req/Sec    35.82k     9.34k   69.40k    70.64%
Latency Distribution
  50%     8.55ms
  75%    14.81ms
  90%    23.92ms
  99%    59.50ms
4195471 requests in 30.08s, 0.88GB read
Requests/sec: 139467.87
Transfer/sec:    29.79MB

```

@greenlaw110 submitted a pull request to add NijinFramework and here is the result.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    55.76ms   63.72ms   1.05s   90.55%
  Req/Sec    12.11k     2.93k   24.49k   77.75%
Latency Distribution
  50%    38.35ms
  75%    71.44ms
  90%   116.40ms
  99%   324.54ms
1442613 requests in 30.08s, 411.36MB read
Requests/sec: 47956.43
Transfer/sec:    13.67MB

```

@greenlaw110 submitted a pull request to add Play1 and here is the result.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:9000 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:9000
4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    10.73ms   23.78ms  397.31ms    0.00%
  Req/Sec    11.32k     2.41k   14.02k    86.33%
Latency Distribution
  50%     0.00us
  75%     0.00us
  90%     0.00us
  99%     0.00us
1335580 requests in 30.02s, 562.98MB read
Requests/sec: 44491.87

```

```
Transfer/sec:      18.75MB
```

@rygel submitted a pull request to add Pippo-Jetty and here is the result.

```
steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
 4 threads and 128 connections
Thread Stats   Avg      Stdev     Max   +/-  Stdev
  Latency    15.66ms   26.96ms  464.86ms   93.37%
  Req/Sec    45.05k    18.29k   100.74k    69.23%
Latency Distribution
 50%     8.11ms
 75%    16.11ms
 90%    31.43ms
 99%   149.57ms
5356385 requests in 30.08s, 807.10MB read
Requests/sec: 178055.45
Transfer/sec:      26.83MB
```

@rygel submitted a pull request to add Pippo-Tomcat and here is the result.

```
steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
 4 threads and 128 connections
Thread Stats   Avg      Stdev     Max   +/-  Stdev
  Latency    23.50ms   40.19ms  768.53ms   93.01%
  Req/Sec    26.57k    11.43k    86.99k    70.86%
Latency Distribution
 50%    12.15ms
 75%    24.84ms
 90%    49.50ms
 99%   204.89ms
3125385 requests in 30.07s, 459.58MB read
Requests/sec: 103948.18
Transfer/sec:      15.29MB
```

@rygel submitted a pull request to add Pippo-Undertow and here is the result.

```
steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
 4 threads and 128 connections
```

```

Thread Stats   Avg      Stdev     Max    +/-  Stdev
Latency       9.80ms   21.25ms  448.49ms  95.90%
Req/Sec       54.74k   16.92k   85.44k    77.53%
Latency Distribution
 50%     5.44ms
 75%     9.18ms
 90%    15.85ms
 99%   109.31ms
6500688 requests in 30.06s, 0.92GB read
Requests/sec: 216254.56
Transfer/sec:   31.35MB

```

@gonzalovazquez submitted a pull request to add node-uws and here is the result.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
Latency       7.22ms   69.98ms   1.64s   99.15%
Req/Sec      148.28k   28.41k   295.52k   93.92%
Latency Distribution
 50%     1.88ms
 75%     2.73ms
 90%     3.24ms
 99%     6.91ms
17706848 requests in 30.02s, 861.21MB read
Socket errors: connect 0, read 0, write 0, timeout 4
Requests/sec: 589924.44
Transfer/sec:   28.69MB

```

@seanjburns submitted a pull request to add vertx and here is the result.

```

wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --latency -- / 16
Running 30s test @ http://localhost:8080
4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
Latency       1.34ms    5.61ms  170.00ms   94.85%
Req/Sec       0.95M   133.67k   1.23M    92.76%
Latency Distribution
 50%    278.00us
 75%    517.00us
 90%     3.23ms
 99%    126.69ms
113247792 requests in 30.10s, 9.70GB read
Requests/sec: 3762376.24
Transfer/sec:   330.10MB

```

@kevin-better submitted a pull request to add baseio and here is the result.

```
wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --latency -- / 16
Running 30s test @ http://localhost:8080
4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    22.20ms   11.17ms  203.34ms   59.61%
  Req/Sec    12.66k     1.85k   15.82k    77.92%
Latency Distribution
  50%    22.16ms
  75%    32.24ms
  90%    37.42ms
  99%    40.12ms
1511579 requests in 30.01s, 191.73MB read
Requests/sec: 50361.98
Transfer/sec: 6.39MB
```

@jknack submitted a pull request to upgrade jooby to version 1.1 and here is the result.

```
Running 30s test @ http://localhost:8080
4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency     4.31ms    3.32ms   48.26ms   78.72%
  Req/Sec    79.88k     7.98k  152.40k   77.13%
Latency Distribution
  50%     3.61ms
  75%     5.74ms
  90%     8.06ms
  99%    16.61ms
9550240 requests in 30.09s, 1.34GB read
Requests/sec: 317385.05
Transfer/sec: 45.70MB
```

@JesHansen submitted a pull request to add dotnet core and here is the result.

```
steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:5000 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:5000
4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency     2.93ms    2.51ms   45.13ms   85.78%
  Req/Sec   122.23k   10.92k  150.99k   72.25%
Latency Distribution
  50%     2.35ms
  75%     3.77ms
```

```

    90%    5.51ms
    99%   12.84ms
14599887 requests in 30.03s, 1.67GB read
Requests/sec: 486216.93
Transfer/sec:    57.03MB

```

@adrienlauer submitted a pull request to add seedstack-filter and here is the result.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
 4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    4.41ms   3.99ms  68.47ms   88.16%
  Req/Sec   86.43k   10.63k  150.72k   74.27%
Latency Distribution
 50%    3.44ms
 75%    5.71ms
 90%    8.79ms
 99%   18.81ms
10336608 requests in 30.10s, 1.51GB read
Requests/sec: 343416.33
Transfer/sec:    51.42MB

```

@adrienlauer submitted a pull request to add seedstack-jersey2 and here is the result.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
 4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency   26.88ms  20.07ms  230.49ms   74.47%
  Req/Sec   13.12k    3.14k   25.28k   75.38%
Latency Distribution
 50%   22.35ms
 75%   36.26ms
 90%   53.21ms
 99%   93.83ms
1574128 requests in 30.09s, 357.29MB read
Requests/sec: 52310.11
Transfer/sec:    11.87MB

```

@jknack submitted a pull request with upgraded Jooby and here is the result.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
 4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency    3.95ms   3.47ms   74.38ms   87.31%
  Req/Sec    91.09k    9.20k   146.91k   75.65%
Latency Distribution
  50%     3.18ms
  75%     5.12ms
  90%     7.52ms
  99%    16.20ms
10896144 requests in 30.10s, 1.41GB read
Requests/sec: 362018.07
Transfer/sec:    47.99MB

```

@rpinaa submitted a pull request to add node-loopback and here is the result.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
 4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency   34.42ms   22.01ms  535.98ms   68.62%
  Req/Sec    8.10k    578.42   11.03k    79.28%
Latency Distribution
  50%    40.16ms
  75%    59.88ms
  90%     0.00us
  99%     0.00us
 965939 requests in 30.10s, 346.37MB read
Requests/sec: 32091.95
Transfer/sec:    11.51MB

```

@rodrigomata submitted a pull request to add go-httprouter and here is the result.

```

steve@joy:~/tool/wrk$ wrk -t4 -c128 -d30s http://localhost:8080 -s pipeline.lua --
latency -- / 16
Running 30s test @ http://localhost:8080
 4 threads and 128 connections
Thread Stats   Avg      Stdev     Max    +/-  Stdev
  Latency   14.12ms   14.25ms  171.89ms   87.11%
  Req/Sec   43.19k    5.18k   61.08k    72.83%
Latency Distribution
  50%     9.58ms
  75%    18.79ms
  90%    32.19ms

```



```
99% 67.61ms
5164540 requests in 30.05s, 635.36MB read
Requests/sec: 171852.31
Transfer/sec: 21.14MB
```