


Software Architecture Patterns

*based on a tutorial of Michael Stal

Harald Gall

University of Zurich

<http://seal.ifi.uzh.ch/ase>



University of
Zurich ^{UZH}



Overview

- Goal
 - Basic architectural understanding of patterns, software architecture and middleware
- Motivation
 - Fundamental architectural principles
 - Platform Comparison
- Design Issues
- Summary

Motivation

- Driven by the Internet as well as mobile and embedded devices **distributed solutions** are now considered common place.
- However, building **distributed applications** is a **non-trivial** task.
- The question is: how can we efficiently build and deploy such applications?

Multi-tier approach

- Separation of concerns
 - Presentation
 - Business logic (middle-tier)
 - Data (backend system)

Multi-tier components

- Presentation Tier components:
 - they typically represent sophisticated GUI elements.
 - they share the same address space with their clients.
 - their clients are containers that provide all the resources.
 - they send events to their containers.
- Middle Tier components:
 - they typically provide server-side functionality.
 - they run in their own address space.
 - they are integrated into a container that hides all system details.

Requirements for Distributed Component-based Applications

- Transparent, platform-neutral communication
- Activation strategies for remote components
- Non-functional properties such as performance, scalability, quality of service (QoS)
- Mechanism to find and create remote components
- Keeping state persistent and consistent
- Security issues
- Data transformation
- Deployment and configuration

Using Standard OO Middleware

- What about using plain CORBA or DCOM?
 - Yes, they provide an OO-RPC
 - But developers must integrate services such as transactions or security themselves
- Integration done by developers may require more than 50% of development time
- Conclusion:
 - An object-oriented RPC helps to connect different islands of code, but is not sufficient for building the middle tier.

Architecture for Middle Tier Components

- Let us introduce step by step some of the architectural elements required to build sophisticated middle tier component infrastructures.
 - Communication


Communication

- In Multi-Tier systems we need to connect the tiers as well as the components within these tiers using communication mechanisms.
- There are four main styles of communication:
 - Collocated client/server (native procedure call)
 - Synchronous RPC
 - Asynchronous RPC
 - Message Queuing (Publish/Subscribe, Peer-to-Peer)

Handling Synchronous Communication

- Components should appear same as local components from a communication perspective
- Clients and servers should be oblivious to communication issues.

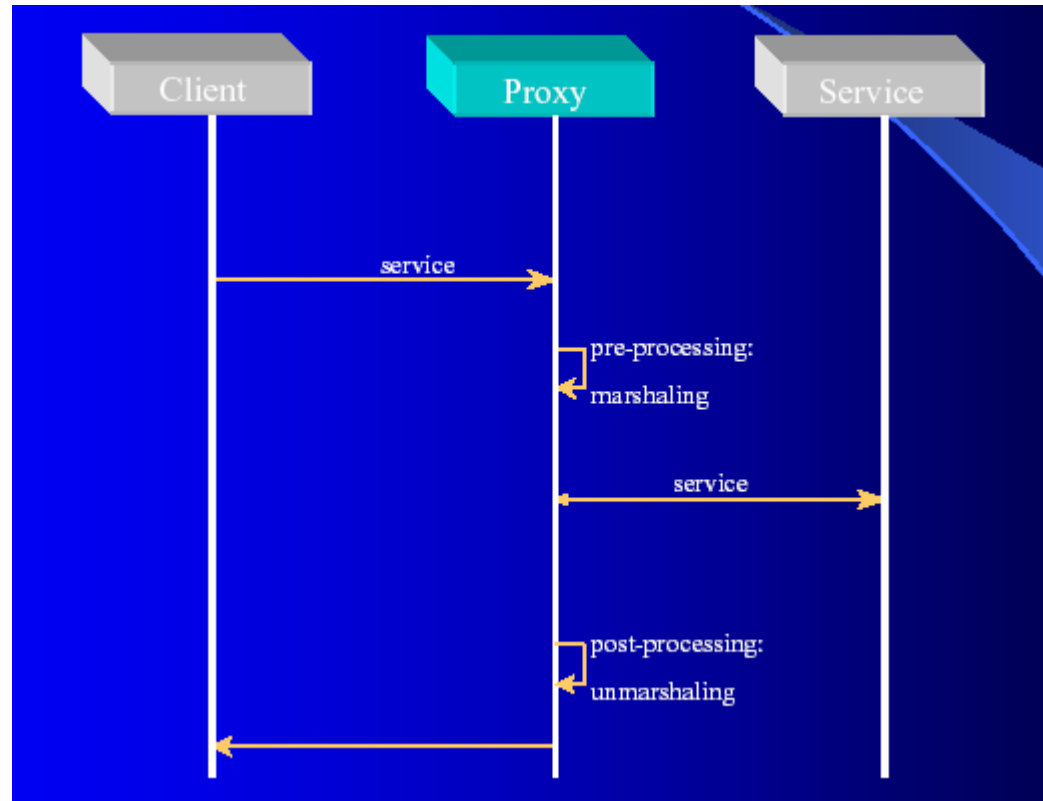
Proxy Pattern



■ Solution

- Provide a placeholder for the object through which clients can access it
- A *Service* implements the object which is not directly accessible.
- A *Proxy* represents the *Service* and ensures the correct access to it. The *Proxy* offers the same interface as the *Service*.
- *Clients* use the *Proxy* to get access to the *Service*.

Dynamics



Benefits and Liabilities

■ Benefits

- Access control to originals.
- Memory savings.
- Performance gaining (cache proxy).
- Separation of housekeeping and functionality.

■ Liabilities

- Potential overkill, if proxies include overly sophisticated functionality.
- Level of indirection.

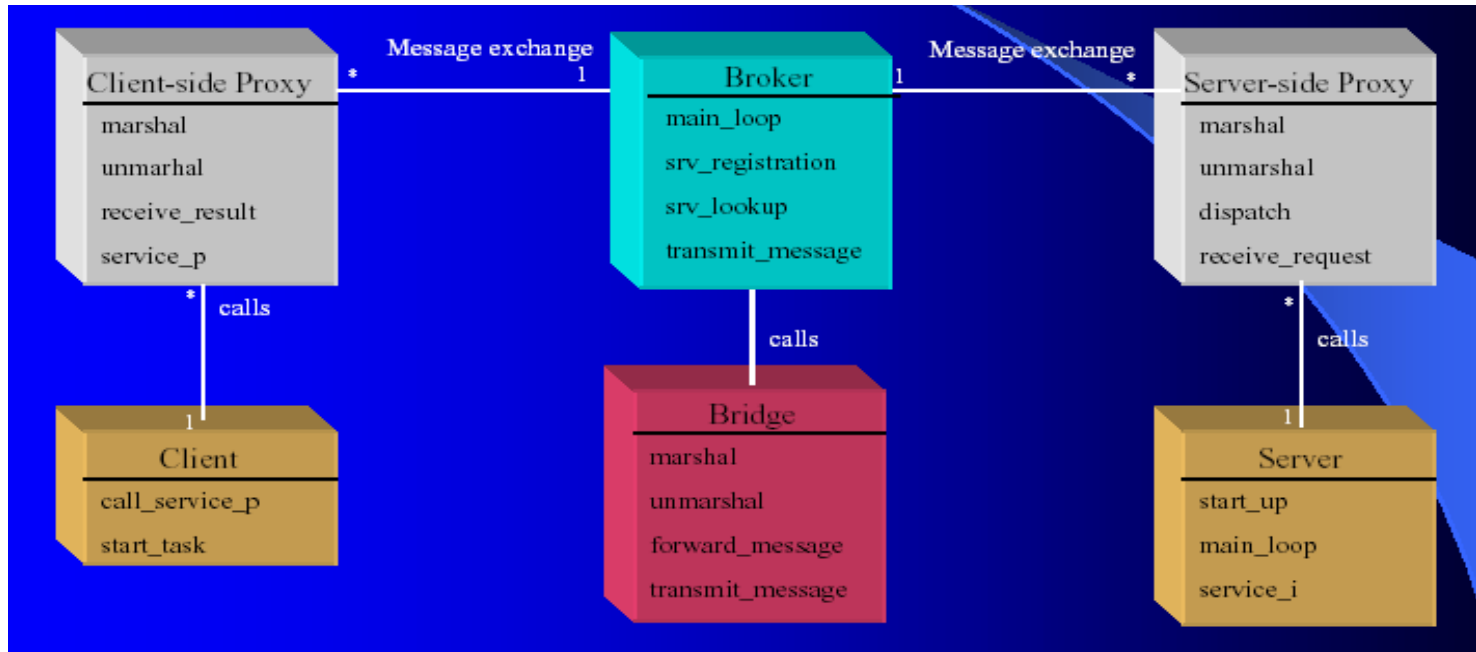
Proxy++

- However, using proxies is not sufficient:
 - How do we locate remote components?
 - How do we handle communication establishment and exchange of network packets (the protocol)?

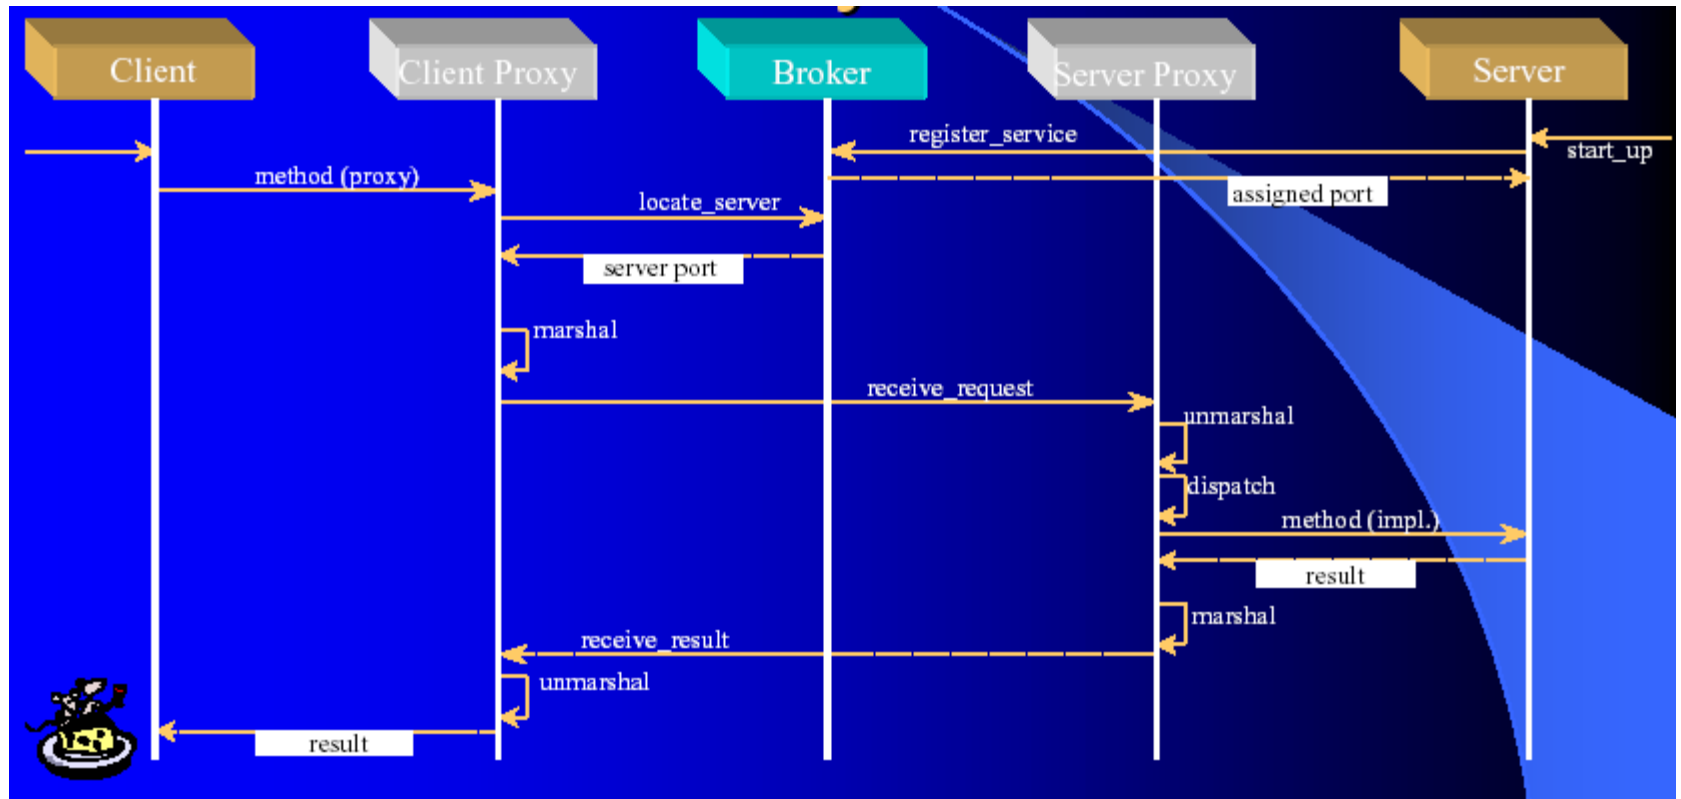
Object-Oriented Middleware Architecture

- What we need is an architecture that ...
 - supports a remote method invocation paradigm
 - provides location transparency
 - allows to add, exchange, or remove services dynamically
 - hides system details from the developer

Architectural Pattern: Broker



Broker Dynamics



Benefits and Liabilities

■ Pros


- Broker and Proxies hide communication
- Details of the OS are also hidden
- Interoperability with other brokers (bridges)
- Reusability of services
- Location Transparency
- Dynamic
- Reconfiguration

■ Cons

- Restricted efficiency due to indirection
- Multiple points of failure
- Testing and debugging harder as with all distributed systems

Generation of „Glue“

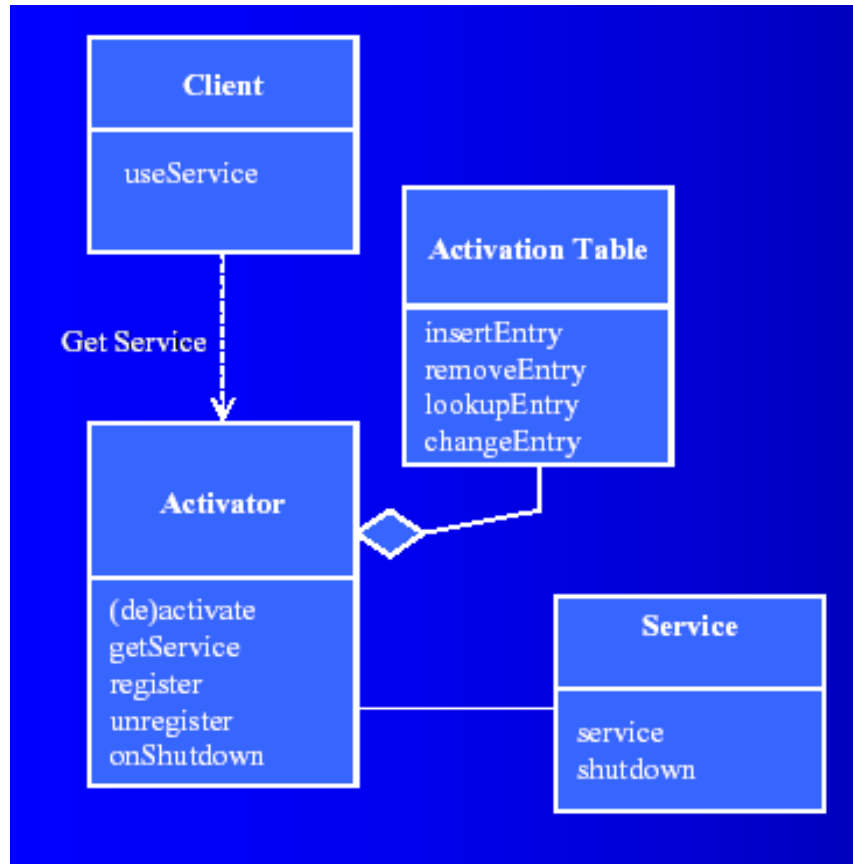
- Who is in charge to implement all of these factories, proxies, etc.?
- Generator tools provide the generation of necessary artifacts from high level language data definitions:



Problem: Activation

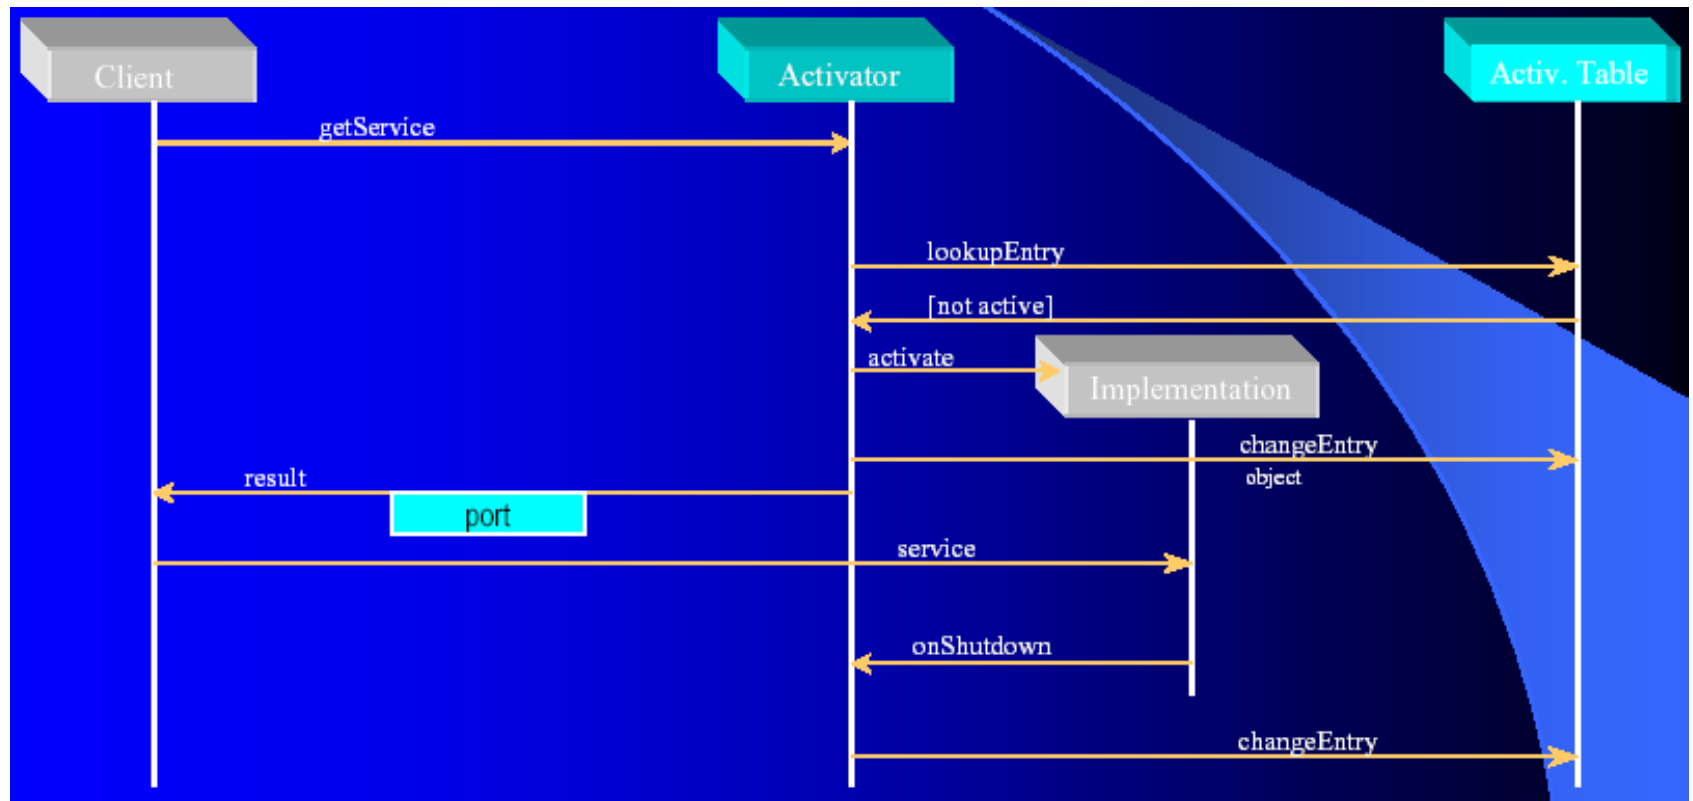
- It is not feasible to have all server implementations running all the time.
- Thus, there must be a way to activate servers on demand.
- Which activation strategy should a broker follow?

Solution



- **Integrate activation code for automatically starting up server implementations. Provide necessary information in tables:**
 - Upon incoming requests the *Activator* looks up whether a target object is already active. If the object is not running it activates the implementation.
 - The *Activation Table* stores associations between services and their physical location.
 - The *Client* uses the *Activator* to get service access.
 - A *Service* implements a specific type of functionality that it provides to clients.

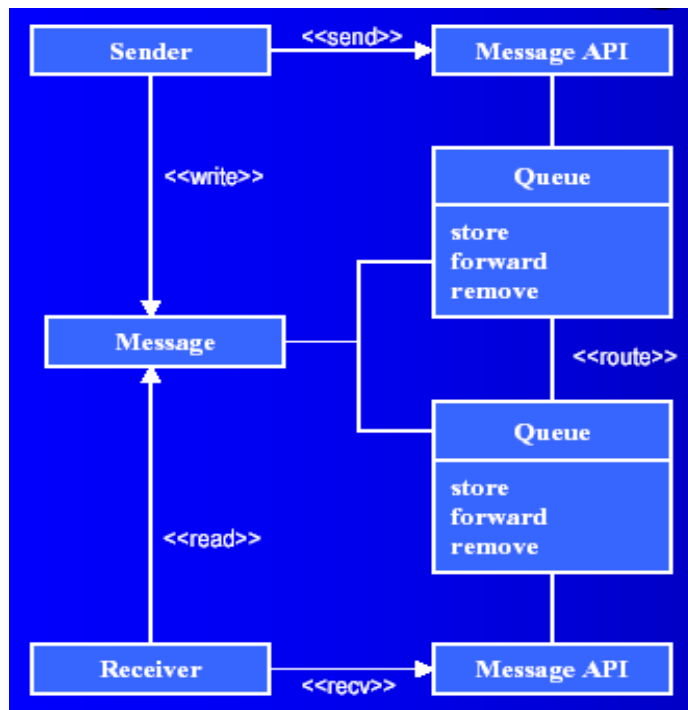
Dynamics: Sample Scenario



Asynchronous and Time Independent Communication

- Using conventional semantics brokers support only blocking communication:
 - Clients must wait until their invocation returns. Even if they use multi-threading the underlying brokers will block.
 - A receiver must be online. However, this cannot always be guaranteed.

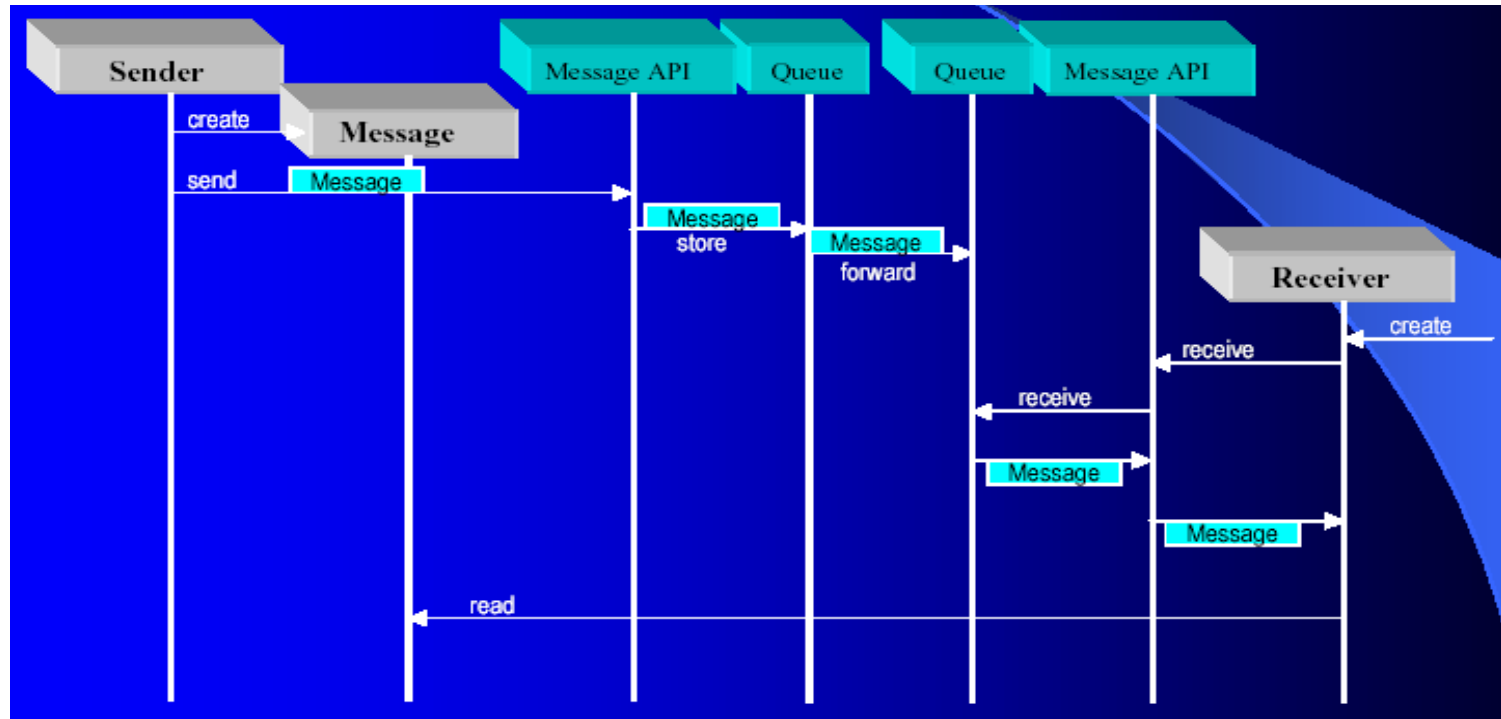
Solution



Solution

- Introduce intermediary queues between senders and receivers :
 - *Queues* are used to store messages persistently.
 - They cooperate with other queues for message routing.
 - *Messages* are objects sent from a sender to a receiver.
 - A *sender* sends messages, while a *receiver* receives them.
 - A *Message API* is provided for senders and receivers to send/receive messages.

Dynamics



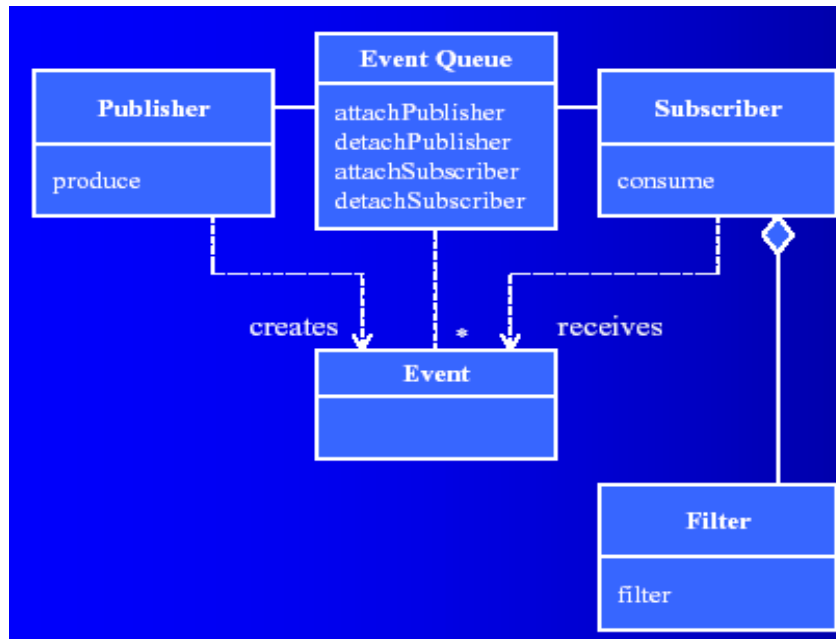
Developer's View

- OO developers prefer request/response semantics:
 - Two message types are used:
 - Requests contain in/in-out arguments
 - Results carry out/inout arguments and results
 - Callback object or Poller object to retrieve result.
 - Player/Recorder strategy: Recorder enqueues request-part of invocation, Player dispatches to actual implementation and enqueues result-part.

Publish/Subscribe

- Usually, a specific client calls a specific remote server and blocks until the result returns.
- Sometimes, this strategy is not sufficient.
 - Consider a server that reports share values.
 - A polling strategy leads to performance bottlenecks.
 - The share values could be spread across different servers.
 - More than one client is interested in the information.
- How can we decouple clients and servers?

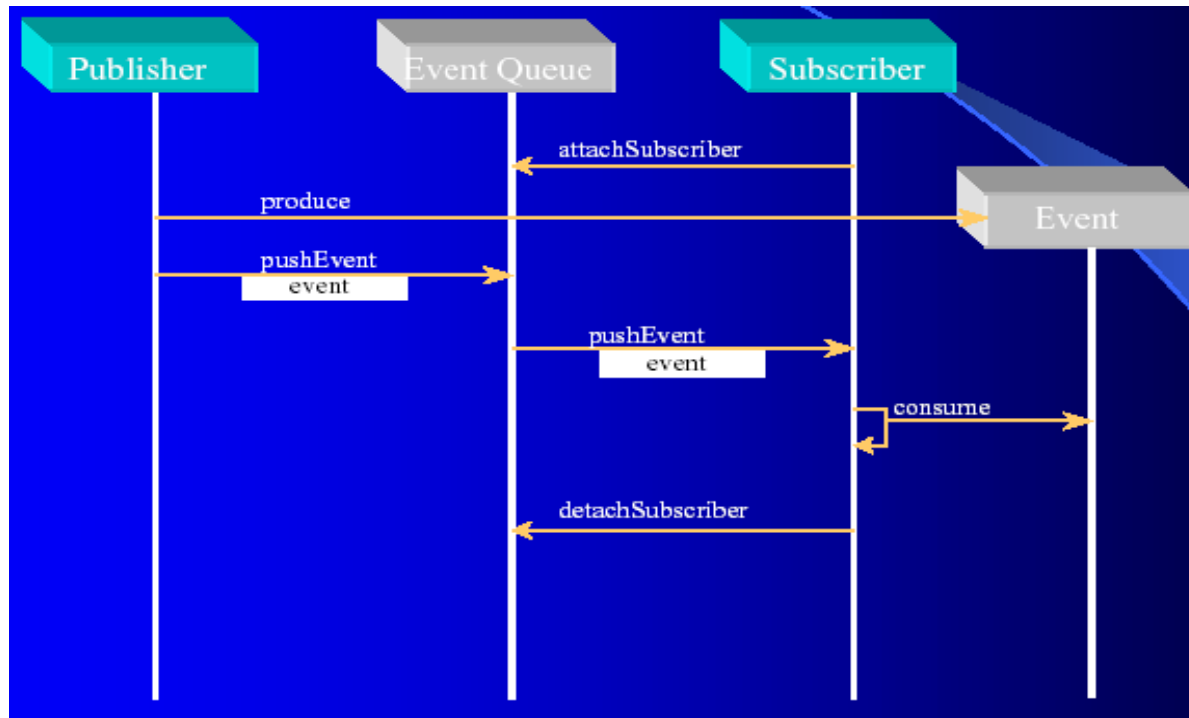
Solution



Decouple suppliers (publishers) and consumers (subscribers) of events:

- An *Event Queue* is storing events.
- *Publishers* create events and store them in an event queue with which they have previously registered.
- *Consumers* register with event queues from which they retrieve events.
- *Events* are objects used to transmit state change information from publishers to consumers.
- For event transmission *push-models* and *pullmodels* are possible.
- *Filters* could be used to filter events on behalf of subscribers.

Dynamics



Benefits and Liabilities

■ Benefits

- Decouples consumers and producers of events.
- n:m communication models are supported.

■ Liabilities

- Must be careful with potential update cascades.

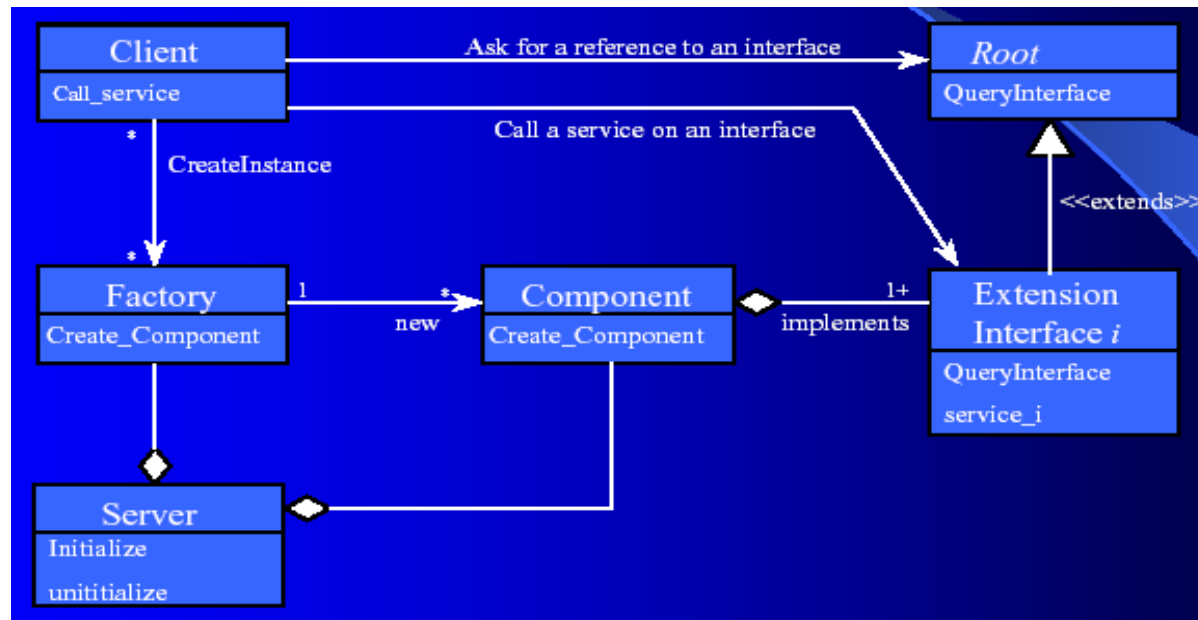
How to provide and evolve component functionality

- Every Component Model defines how components should import and export functionality.
- Components should be able to evolve over time without impact on clients.

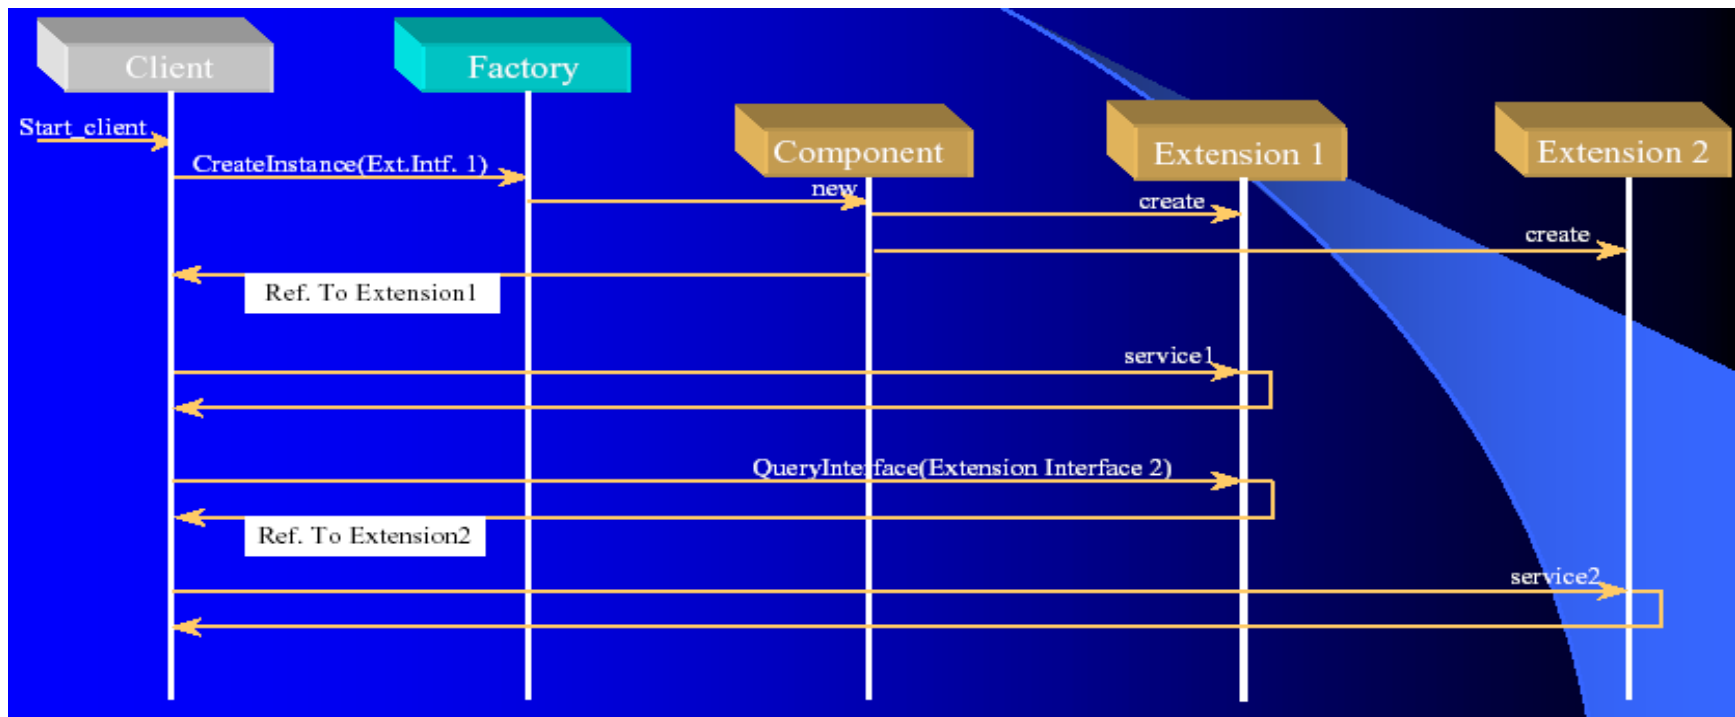
Multiple Interfaces

- Most objects have only one interface.
- Thus, if we want the banana, we get the whole gorilla.
- How can we make this more flexible?

Extension Interface Pattern



Dynamics



Benefits and Liabilities

■ Benefits

- Exchangeability of components
- Extensibility through interfaces
- Prevention of interface bloating
- No subclassing required
- Separation of concerns

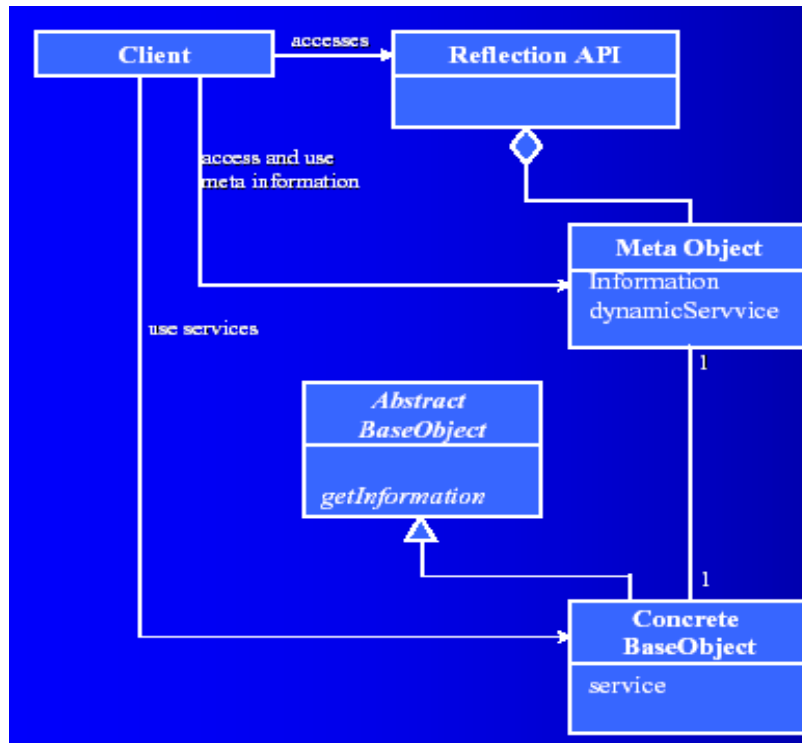
■ Liabilities

- Restricted efficiency due to indirection
- Complexity and cost for development and deployment

Dynamic Environments

- Until now, clients can only access services that are available at compile-time.
- Typically, in distributed environments services must be dynamically added, exchanged, or removed.
- How can we access objects that are not known at deployment-time?

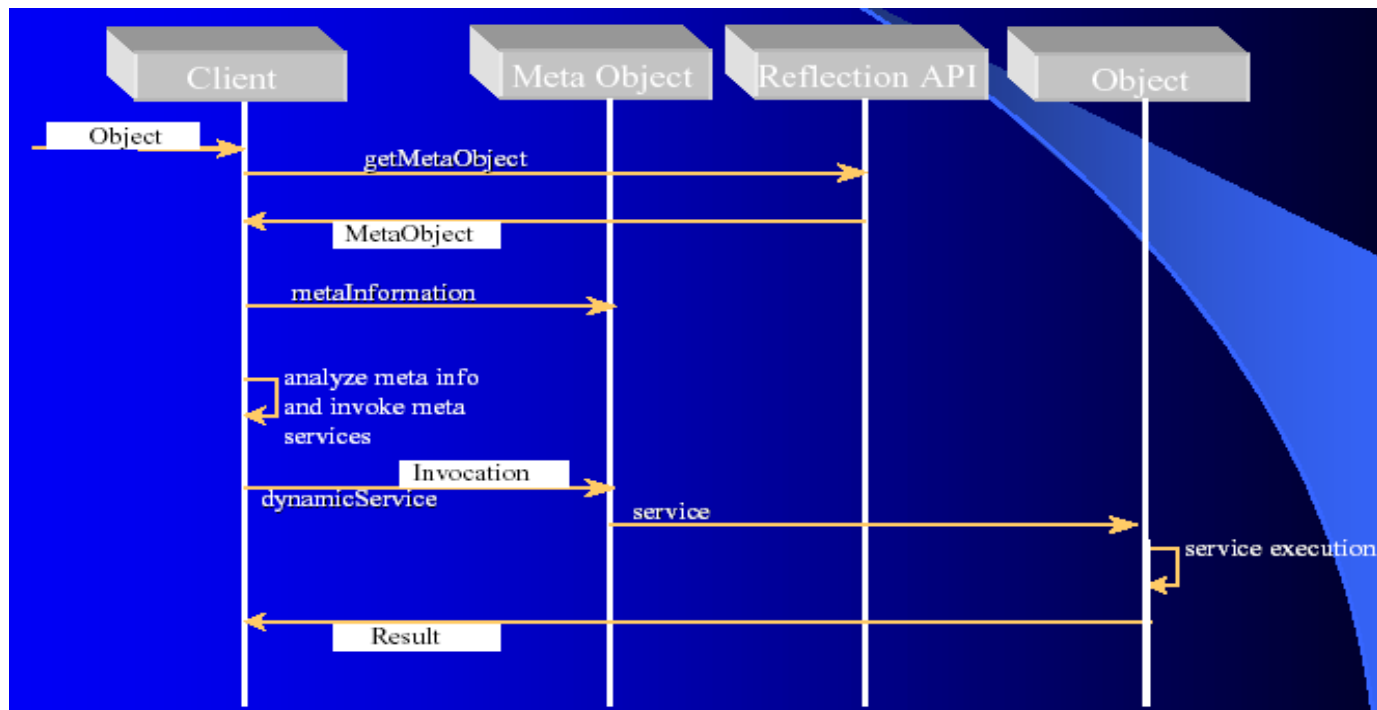
Simplified Reflection Pattern



Solution

- Provide type information through meta objects and allow clients to find and use this meta information:
 - *Meta Objects* provide information and functionality about existing base objects.
 - The *Reflection API* is responsible for creating and retrieving Reflection Objects.
 - *Base Objects* implement the application functionality.
 - *Clients* access Meta Objects to dynamically invoke base functionality.

Dynamics



Benefits and Liabilities

■ Benefits

- Support for change.
- Clients may dynamically invoke objects without compile-time knowledge.
- Support for visual builder tools and for scripting environments.
- Base Objects may be added at run-time.

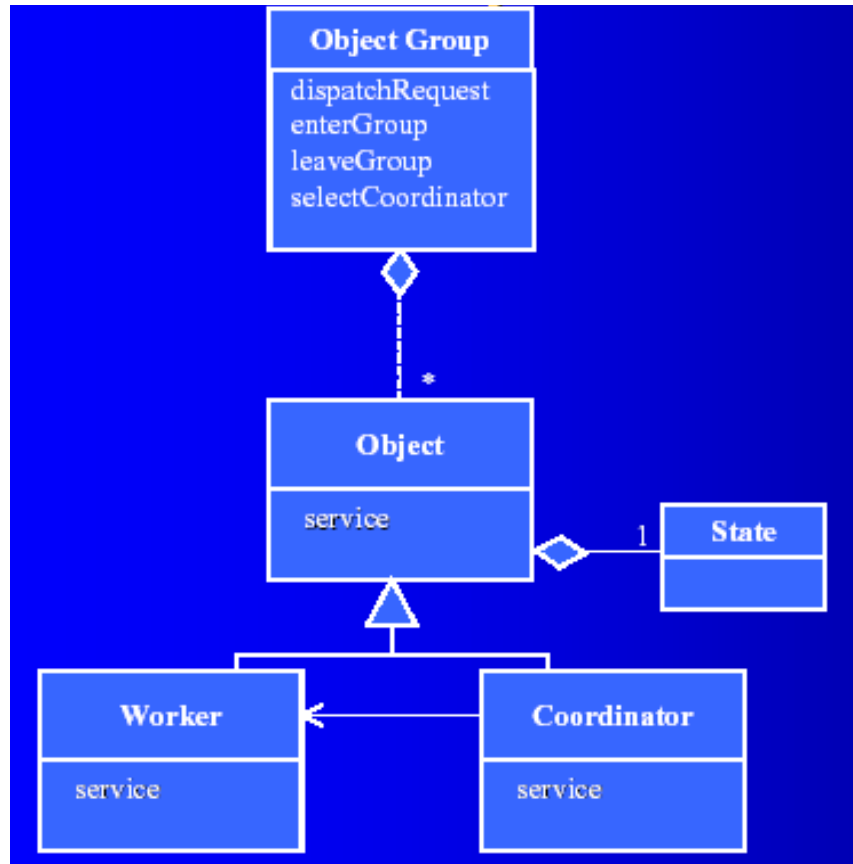
■ Liabilities

- Possible performance problems.
- Implementation harder to maintain and understand.

Fault Tolerance

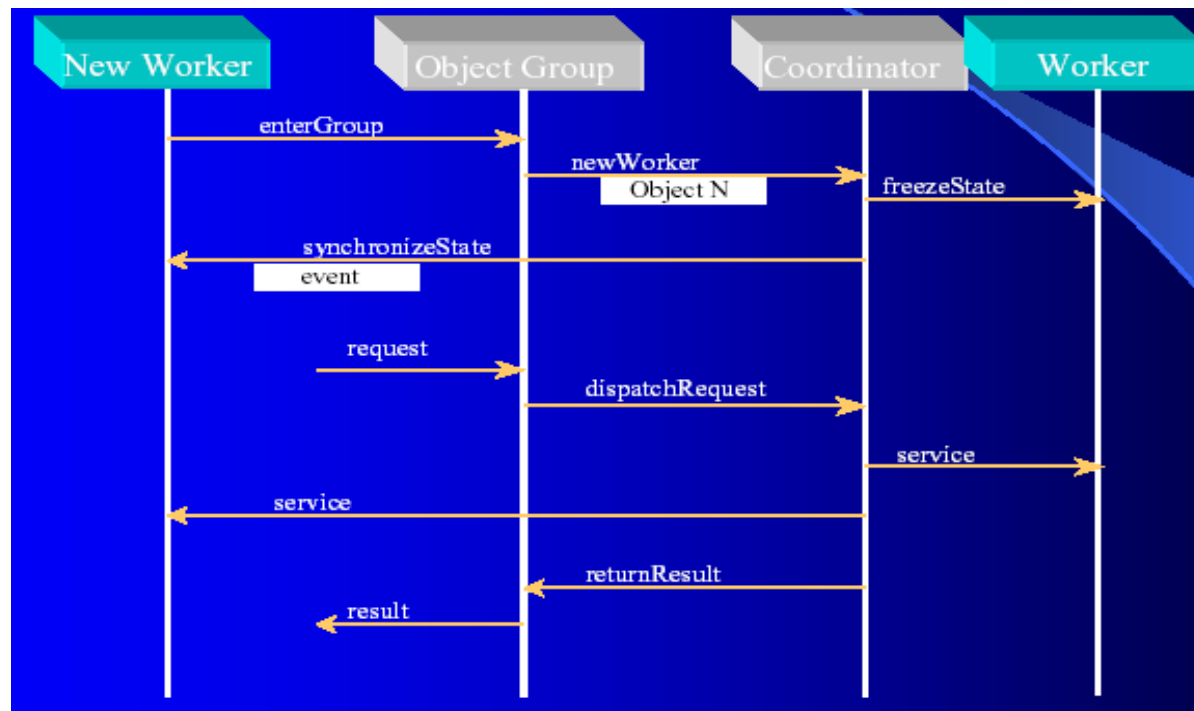
- To enhance fault tolerance, more than one server should provide the same service.
- It should be transparent to the client that replicated servers are used.
- How can we introduce this kind of transparency?

Object Group Pattern



- **Decouple suppliers (publishers) and consumers (subscribers) of events:**
 - An Object Group represents a group of objects that all provide the same service.
 - Object provides a specific service to clients.
 - The coordinator is selected by the Object Group to dispatch requests.
 - Workers receive requests from their coordinator and return results.
 - Each member of the object group maintains its own state.

Dynamics



Benefits and Liabilities

■ Benefits

- Resilience to failures in distributed object implementations.
- Clients unaffected by object group internals.
- State remains consistent.

■ Liabilities

- Overhead due to several indirection layers.

Conclusions

- Software architecture patterns as
 - a blueprint for architectural solutions
 - guidance for benefits and liabilities
 - for architecture assessment